

Charles Hong

charleshong@berkeley.edu | charleshong3.github.io

EDUCATION

University of California, Berkeley <i>Ph.D., Electrical Engineering and Computer Science</i>	Ongoing
University of California, Berkeley <i>M.S., Electrical Engineering and Computer Science</i>	2025
University of California, Berkeley <i>B.S., Electrical Engineering and Computer Science</i>	2021

HONORS/AWARDS

LAD Fellowship – IEEE	2025
Master's Scholarship in Integrated Systems – Apple	2021
High Honors at Graduation – University of California, Berkeley	2021
Eta Kappa Nu Honor Society – IEEE	2018
Edward Frank Kraft Award for Freshmen – University of California, Berkeley	2017

CONFERENCE PAPERS

- hdl2v: A Code Translation Dataset for Enhanced LLM Verilog Generation** ([Paper](#)) ([Dataset](#))
Charles Hong, Brendan Roberts, Huijia An, Alex Um, Advay Ratan, Yakun Sophia Shao
7th ACM/IEEE International Symposium on Machine Learning for CAD (MLCAD), September 2025
- DOSA: Differentiable Model-Based One-Loop Search for DNN Accelerators** ([Paper](#)) ([Slides](#)) ([GitHub](#))
Charles Hong, Qijing Huang, Grace Dinh, Mahesh Subedar, Yakun Sophia Shao
56th IEEE/ACM International Symposium on Microarchitecture (MICRO), October 2023
- Learning A Continuous and Reconstructible Latent Space for Hardware Accelerator Design** ([Paper](#)) ([Slides](#)) ([GitHub](#))
Qijing Huang, Charles Hong, John Wawrzynek, Mahesh Subedar, Yakun Sophia Shao
International Symposium on Performance Analysis of Systems and Software (ISPASS), May 2022

PREPRINTS

- Autocomp: A Powerful and Portable Code Optimizer for Tensor Accelerators** ([Paper](#)) ([Blog](#)) ([GitHub](#))
Charles Hong, Sahil Bhatia, Alvin Cheung, Yakun Sophia Shao
arXiv, May 2025 (under submission)
- Polaris: Multi-Fidelity Design Space Exploration of Deep Learning Accelerators** ([Paper](#))
Chirag Sakhuja, Charles Hong, Calvin Lin
arXiv, December 2024

WORKSHOPS AND RESEARCH COMPETITIONS

- Autocomp: LLM-Driven Code Optimization for Tensor Accelerators** ([Paper/Talk/Slides](#))
Outstanding Paper Award
Charles Hong, Sahil Bhatia, Alvin Cheung, Yakun Sophia Shao
ML for Computer Architecture and Systems Workshop (MLArchSys), co-located with ISCA, June 2025
- LLM-Aided Compilation for Tensor Accelerators** ([Paper](#))
Charles Hong, Sahil Bhatia, Altan Haan, Shengjun Kris Dong, Dima Nikiforov, Alvin Cheung, Yakun Sophia Shao
The First IEEE International Workshop on LLM-Aided Design (ISLAD), June 2024

Sample-Efficient Mapspace Optimization for DNN Accelerators with Bayesian Learning (Paper)

Grace Dinh, Iniyaal Kannan, Hengrui Luo, Charles Hong, Younghyun Cho, James Demmel, Sherry Li, Yang Liu
ML for Computer Architecture and Systems Workshop (MLArchSys), co-located with ISCA, June 2023

Modeling DNN Layer Performance Across Accelerators (Paper) (Poster)

Charles Hong, Yakun Sophia Shao
MICRO ACM Student Research Competition, October 2021

EXPERIENCE

University of California, Berkeley

June 2020 – Present

Undergraduate/Graduate Student Researcher

Berkeley, CA

- Researcher at ADEPT and SLICE computer architecture labs, advised by Professor Sophia Shao.
- Areas of focus include deep learning accelerators, hardware/software co-design, and ML for systems (see publications and projects for details).

Siemens EDA

May 2024 - August 2024

HLS Intern

Wilsonville, OR

- Improved tool flow for Catapult HLS-generated machine learning acceleration and demonstrated results via FPGA prototype-based demo.

University of California, Berkeley

August 2018 – August 2023

Undergraduate/Graduate Student Instructor

Berkeley, CA

- Lecturer for CS 61C in Summer 2023.
- Six-time (U)GSI for CS 61A (programming fundamentals), CS 61C (computer architecture), and EECS 151 (digital design, FPGAs).

Intel Labs

May 2022 - December 2022

AI Research Intern

Remote

- Developed methods for AI/ML-based optimization of deep learning accelerator hardware and software.
- Prototyped an Intel FPGA-based platform for SoC evaluation.

Apple

May 2021 – August 2021

Hardware Technology Intern

Cupertino, CA

- Intern in CPU DV. Developed CPU Emulation Flow 2.0, which enhances support for internal emulation users, increases test throughput, and reduces runtime of common tasks.

NVIDIA

May 2020 – August 2020

System Architect Intern

Santa Clara, CA

- Profiled CPU networking workloads and built a machine learning-based cost model, predicting cycle counts to within 15% accuracy.

NVIDIA

May 2019 – August 2019

System Architect Intern

Santa Clara, CA

- Built tool to model power data of SoC workloads based on hardware configuration and task scheduling across various IPs.
- Enabled estimation of power variation over time and sped up evaluation by several times.

Oski Technology

June 2018 – August 2018

Formal Verification Intern

San Jose, CA

- Verified an interconnect bridge design and prototyped Chisel FV flow.

TECHNICAL SKILLS

Software Engineering: Python, Java, C/C++

Web Development: HTML/CSS/JavaScript

Hardware and Systems Engineering: Expert in computer architecture, especially domain-specific accelerators. Experienced in Verilog/SystemVerilog RTL design. Familiar with compilers, egraphs, code generation.

Machine Learning/Data Science: Experienced in model training from sklearn to PyTorch to HuggingFace Transformers. Experienced in data analysis and visualization with NumPy, Pandas, Matplotlib.

EDA Tools: FPGAs (Xilinx Vivado and Altera Quartus), Cadence Joules

Formal Methods: Familiar with Rosette and Z3 for program synthesis and verification, as well as JasperGold for RTL FV.

VOLUNTEERING

Pioneers in Engineering

September 2017 – Present

Advisor, Director of Engineering, Project Manager, Software Developer

Berkeley, CA

- Various roles helping run robotics competition for local underserved high schools.
- Two years as project manager; one year as engineering director, managing 30+.
- Developed robot simulator web app (pimulator.pierobotics.org).

OTHER PROJECTS

Rewrite Rule Synthesis: A Survey ([Paper](#)/[Slides](#))

Charles Hong

CS 294-260 (Declarative Program Analysis and Optimization) Class Project, Spring 2024

Transformers in Reinforcement Learning for Hardware Design Space Exploration ([Paper](#))

Junsun Choi, Charles Hong

CS 285 (Deep Reinforcement Learning) Class Project, Fall 2022

Fast Thread Migration in a Heterogeneous ISA System ([Paper](#)) ([Poster](#))

Max Banister, John Fang, Charles Hong

CS 262A (Advanced Computer Systems) Class Project, Fall 2021

Enhancing Yelp Rating Predictions ([Paper](#))

Charles Hong, Kevin Zhu

CS 182 (Deep Learning) Class Project, Spring 2021

A Chipyard Comparison of NVDLA and Gemmini ([Paper](#))

Abraham Gonzalez, Charles Hong

EE 290-2 (Hardware for Machine Learning) Class Project, Spring 2020